

```
In [25]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import numpy as np
```

```
In [12]: df = pd.read_csv('adult.csv', header=None)
```

```
In [13]: df.head(3)
```

Out[13]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40

```
In [14]: columns = ['age', 'workclass', 'fnlwgt', 'education', 'education-num',
'marital-status',
'occupation', 'relationship', 'race', 'sex', 'capital-gain', 'capital-loss',
'hours-per-week', 'native-country', 'target']

df.columns = columns
```

```
In [15]: df.head(3)
```

Out[15]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male

In [16]: `df.tail(3)`

Out[16]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race
32558	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White
32559	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White
32560	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White

In [17]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         32561 non-null  int64
1   workclass   32561 non-null  object
2   fnlwgt      32561 non-null  int64
```

```

3    education      32561 non-null object
4    education-num  32561 non-null int64
5    marital-status 32561 non-null object
6    occupation     32561 non-null object
7    relationship   32561 non-null object
8    race           32561 non-null object
9    sex            32561 non-null object
10   capital-gain   32561 non-null int64
11   capital-loss   32561 non-null int64
12   hours-per-week 32561 non-null int64
13   native-country 32561 non-null object
14   target         32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```
In [18]: df[['education', 'education-num']].head()
```

Out[18]:

	education	education-num
0	Bachelors	13
1	Bachelors	13
2	HS-grad	9
3	11th	7
4	Bachelors	13

```
In [19]: target = df['target']
target.head(3)
```

Out[19]:

```

0    <=50K
1    <=50K
2    <=50K
Name: target, dtype: object

```

```
In [20]: cat = ['workclass', 'education', 'marital-status', 'occupation', 'relationship',
               'race', 'sex', 'native-country']
```

```
In [21]: df_cat = df[cat]
df_cat.head(3)
```

Out[21]:

	workclass	education	marital-status	occupation	relationship	race	sex	native-country
0	State-gov	Bachelors	Never-married	Adm-clerical	Not-in-family	White	Male	United-States
1	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	United-States
2	Private	HS-grad	Divorced	Handlers-cleaners	Not-in-family	White	Male	United-States

```
In [22]: num = [i for i in df.columns if i not in cat]
_ = df[num]
df_num = _.drop('target', axis=1).drop('education-num', axis=1)
```

```
In [23]: df_num.head(3)
```

Out[23]:

	age	fnlwgt	capital-gain	capital-loss	hours-per-week
0	39	77516	2174	0	40
1	50	83311	0	0	13
2	38	215646	0	0	40

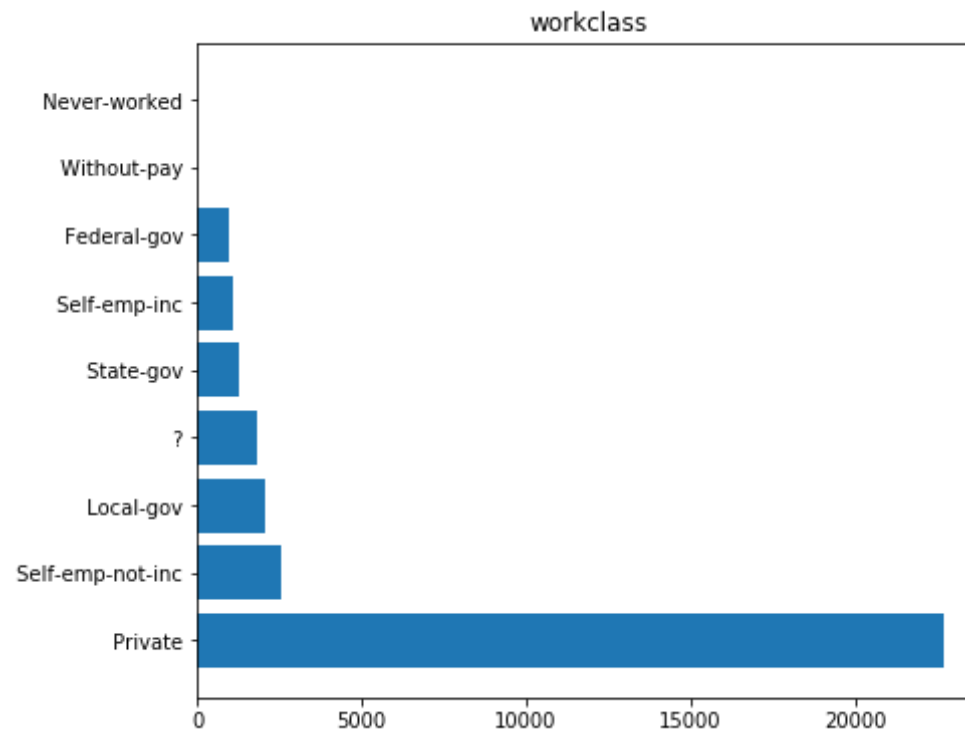
```
In [26]: for n in range(len(df_cat.columns)):
plt.figure(figsize=(7, 6))

keys = list(dict(df_cat.iloc[:, n].value_counts()).keys())
vals = list(dict(df_cat.iloc[:, n].value_counts()).values())

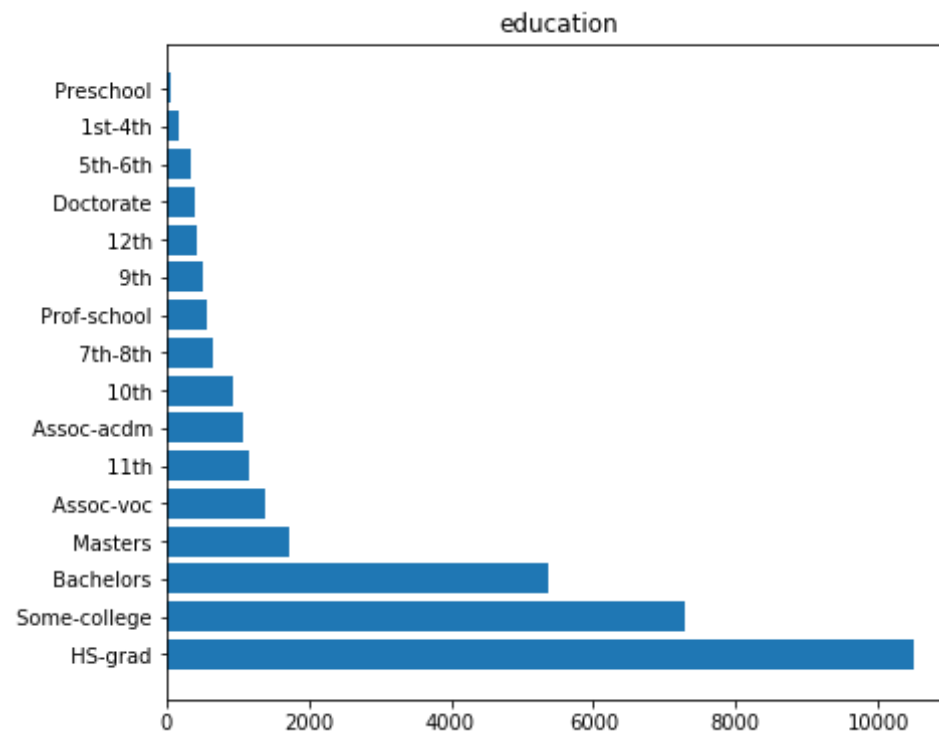
plt.barh(keys, vals)
plt.title(df_cat.columns[n])
```

```
plt.tight_layout
plt.show()

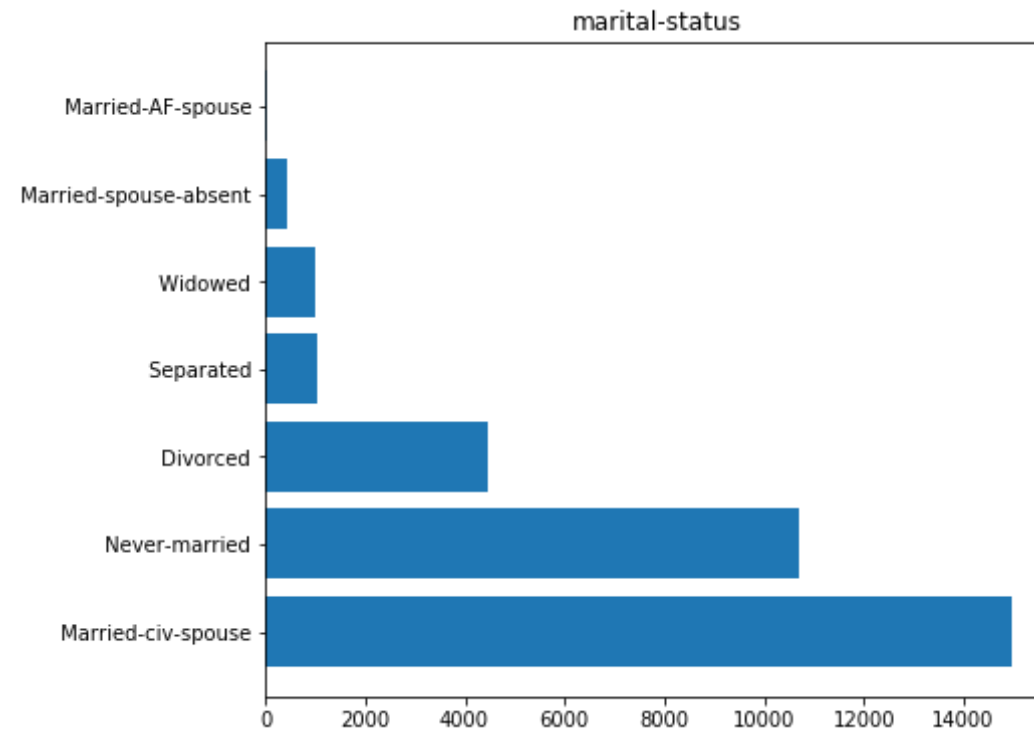
print('{0:.2f}%'.format(max(vals)/np.sum(vals)*100), keys[vals.index(max(vals))])
print()
```




69.70% Private




32.25% HS-grad




45.99% Married-civ-spouse







12.71% Prof-specialty









40.52% Husband







85.43% White



66.92% Male





89.59% United-States


```
In [27]: for n in range(len(df_num.columns)):
plt.figure(figsize=(7, 6))

vals = list(df_num.iloc[:, n])

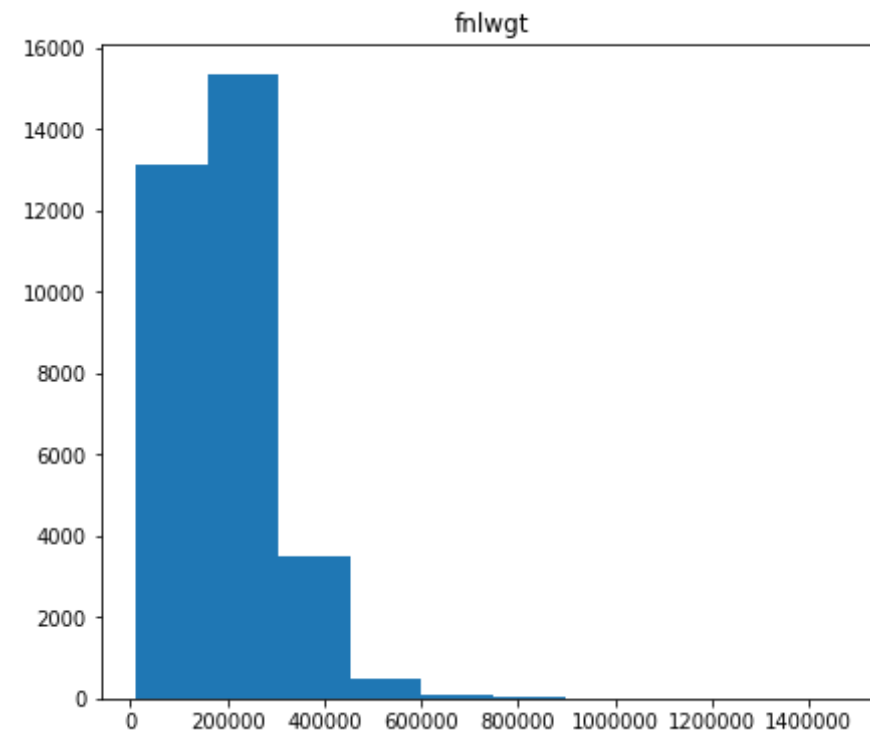
plt.hist(vals)
plt.title(df_num.columns[n])

plt.tight_layout
plt.show()

print('Mean: {0:.2f}'.format(np.mean(vals)))
```




Mean: 38.58

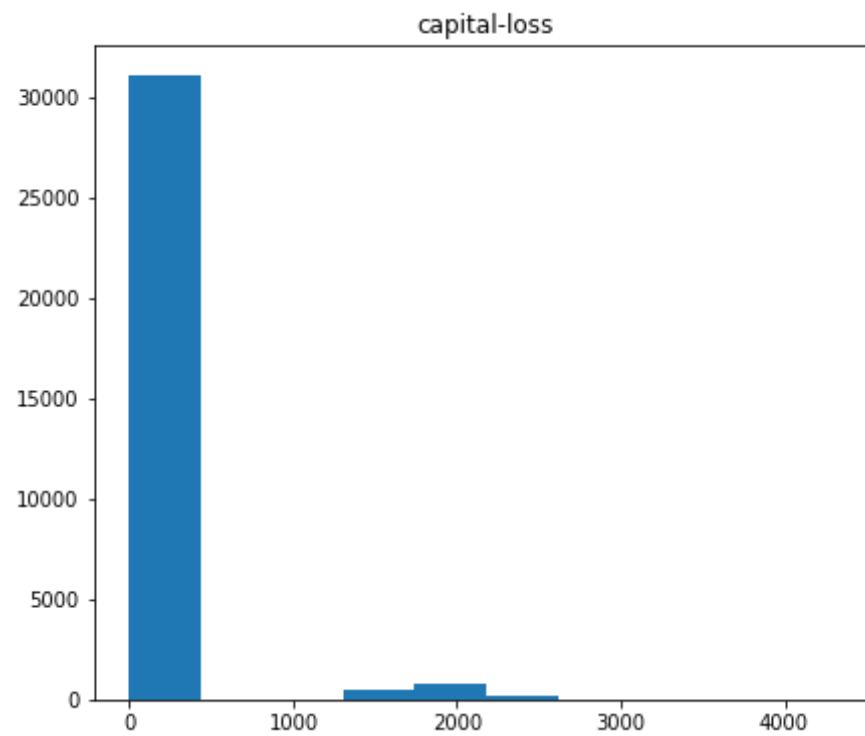


Mean: 189778.37






Mean: 1077.65



Mean: 87.30




Mean: 40.44

```
In [28]: target.value_counts()
```

```
Out[28]: <=50K    24720  
>50K       7841  
Name: target, dtype: int64
```

```
In [29]: _ = dict(target.value_counts())  
plt.bar(_.keys(),_.values())
```

```
Out[29]: <BarContainer object of 2 artists>
```



```
In [30]: df_cat2 = pd.get_dummies(df_cat)
```

```
In [31]: df_cat2.shape
```


```
Out[31]: (32561, 102)
```

```
In [32]: df_cat2.head(3)
```

```
Out[32]:
```

	workclass_?	workclass_Federal-gov	workclass_Local-gov	workclass_Never-worked	workclass_Private	workclass_Self-emp-inc	workclass_Self-emp-not-inc	workclass_State-gov
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	1	0
2	0	0	0	0	1	0	0	0

3 rows × 102 columns




```
In [34]: df_cat2.tail(3)
```



Out[34]:

	workclass_ ?	workclass_ Federal- gov	workclass_ Local-gov	workclass_ Never- worked	workclass_ Private	workclass_ Self-emp- inc	workclass_ Self-emp- not-inc	wo S
32558	0	0	0	0	1	0	0	
32559	0	0	0	0	1	0	0	
32560	0	0	0	0	0	1	0	

3 rows × 102 columns



```
In [35]: _ = pd.get_dummies(pd.DataFrame(target))
```

```
In [36]: target_num = _.iloc[:, 1]
target_num.head()
```

```
Out[36]: 0    0
1    0
2    0
3    0
4    0
Name: target_ >50K, dtype: uint8
```

```
In [37]: target.tail()
```

```
Out[37]: 32556    <=50K
32557    >50K
32558    <=50K
32559    <=50K
32560    >50K
Name: target, dtype: object
```

```
In [38]: target_num.tail()
```

```
Out[38]: 32556    0
32557    1
32558    0
```

```
32559    0
32560    1
Name: target_>50K, dtype: uint8
```

```
In [39]: target_num.shape
```

```
Out[39]: (32561,)
```

```
In [40]: df_num.shape
```

```
Out[40]: (32561, 5)
```

```
In [41]: df_cat2.shape
```

```
Out[41]: (32561, 102)
```

```
In [42]: df_final = pd.concat([df_num, df_cat2, target_num], axis=1)
```

```
In [43]: df_final.shape
```


```
Out[43]: (32561, 108)
```

```
In [44]: df_final.head(3)
```

```
Out[44]:
```

	age	fnlwgt	capital- gain	capital- loss	hours- per- week	workclass_ ?	workclass_ Federal- gov	workclass_ Local-gov	workclass_ Never- worked	workclass_ F
0	39	77516	2174	0	40	0	0	0	0	
1	50	83311	0	0	13	0	0	0	0	
2	38	215646	0	0	40	0	0	0	0	

3 rows × 108 columns



```

In [46]: from collections import Counter

from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier
from sklearn.neighbors import KNeighborsRegressor

from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso

import mglearn

In [47]: def ml_model(feature, target, ml_type='knn_class', show_PCC=False,
                    test_size=0.25, param_range=range(1, 30),
                    seed_settings=range(0, 30), max_iter = 10000, plot = False
                    ,
                    report = True):
    """
    Plot accuracy vs parameter for test and training data. Print
    maximum accuracy and corresponding k value. Print number of trials.

    Inputs
    =====
    target: Dataframe or Series of target variables
    feature: Dataframe of features
    test_size: Float indicating the proportion of data for testing
    target: Series of target values
    test_size: List of values for percentage of test size
    ml_type: String indicating which ML algorithm to use
    param_range: Range of values for parameters

    Outputs
    =====
    Plot of accuracy vs parameter for test and training data
    """

    train_acc = []

```

```

test_acc = []

lasso_feats = []

# Initiate counter for number of trials
iterations = 0

# create an array of cols: parameters and rows: seeds
for seed in seed_settings:

    # count one trial
    iterations += 1

    # split data into test and training sets
    X_train, X_test, y_train, y_test = train_test_split(feature,
                                                         target,
                                                         test_size=t
est_size,
                                                         random_stat
e=seed)
    train = []
    test = []
    lasso = []

    # make a list of accuracies for different parameters
    for param in param_range:
        # build the model
        if ml_type == 'knn_class':
            clf = KNeighborsClassifier(n_neighbors=param)

        elif ml_type == 'knn_reg':
            clf = KNeighborsRegressor(n_neighbors=param)

        elif ml_type == 'ridge':
            clf = Ridge(alpha=param)

        elif ml_type == 'lasso':
            clf = Lasso(alpha=param, max_iter=max_iter)

```

```

        clf.fit(X_train, y_train)
        # record training set accuracy
        train.append(clf.score(X_train, y_train))
        # record generalization accuracy
        test.append(clf.score(X_test, y_test))

    if ml_type == 'lasso':
        lasso.append(np.sum(clf.coef_ != 0))

    # append the list to _acc arrays
    train_acc.append(train)
    test_acc.append(test)

    if ml_type == 'lasso':
        lasso_feats.append(lasso)

# compute mean and error across columns
train_all = np.mean(train_acc, axis=0)
test_all = np.mean(test_acc, axis=0)

# compute standard deviation
std_train = np.std(train_acc, axis=0)
std_test = np.std(test_acc, axis=0)

# compute pcc
state_counts = Counter(target)
df_state = pd.DataFrame.from_dict(state_counts, orient='index')
num = (df_state[0] / df_state[0].sum())**2
pcc = 1.25 * num.sum()
if plot == True:
    plt.figure(figsize=(8, 6))

    # plot train and errors and standard devs
    plt.plot(param_range, train_all, c='b', label="training set", marker='.')
    plt.fill_between(param_range,
                     train_all + std_train,
                     train_all - std_train,
                     color='b', alpha=0.1)

```

```

# plot test and errors and standard devs
plt.plot(param_range, test_all, c='r', label="test set", marker
='.')
plt.fill_between(param_range,
                  test_all + std_test,
                  test_all - std_test,
                  color='r', alpha=0.1)

# plot pcc line
if show_PCC == True:
    plt.plot(param_range, [pcc] * len(param_range), c = 'tab:gr
ay', label = "pcc", linestyle = '--')

plt.xlabel('Parameter Value')
plt.ylabel('Accuracy')
plt.title(ml_type + ": Accuracy vs Parameter Value")
plt.legend(loc = 0)

plt.tight_layout()
plt.show()

max_inds=np.argsort(test_all)[-1]
acc_max=test_all[max_inds]
param_max = list(param_range)[max_inds]
if report == True:
    print('Report:')
    print('=====')
    print("Max accuracy: {}\nOptimal parameter: {}".format(
        np.round(acc_max, 4), param_max))
    if ml_type == 'lasso':
        lasso_feats_mean = np.mean(lasso_feats, axis=0)[max_inds]
        print('Ave no. of features for max accuracy: {} out of {}'.
format(lasso_feats_mean,
len(X[
0]))))
    print('1.25 x PCC: {:.4f}'.format(pcc))
    print('Total iterations: {}'.format(iterations))

```


```
# return lasso_feats
return np.round(acc_max, 4), param_max
```

```
In [48]: X = df_final.iloc[:, :-1]
X.head(3)
```

Out[48]:

	age	fnlwgt	capital- gain	capital- loss	hours- per- week	workclass_ ?	workclass_ Federal- gov	workclass_ Local-gov	workclass_ Never- worked	workclass_ F
0	39	77516	2174	0	40	0	0	0	0	
1	50	83311	0	0	13	0	0	0	0	
2	38	215646	0	0	40	0	0	0	0	

3 rows × 107 columns



```
In [49]: y = target_num
y.head(3)
```

Out[49]:

0	0
1	0
2	0

Name: target\_ >50K, dtype: uint8

```
In [50]: y.tail(3)
```

Out[50]:

32558	0
32559	0
32560	1


Name: target\_ >50K, dtype: uint8

```
In [51]: state_counts = Counter(y)
df_state = pd.DataFrame.from_dict(state_counts, orient='index')
df_state.plot(kind='bar')
```

```
num=(df_state[0]/df_state[0].sum())**2
print("Population per class: {}".format(df_state))
print("1.25 * Proportion Chance Criterion: {}".format(1.25*100*num.sum
()))
```

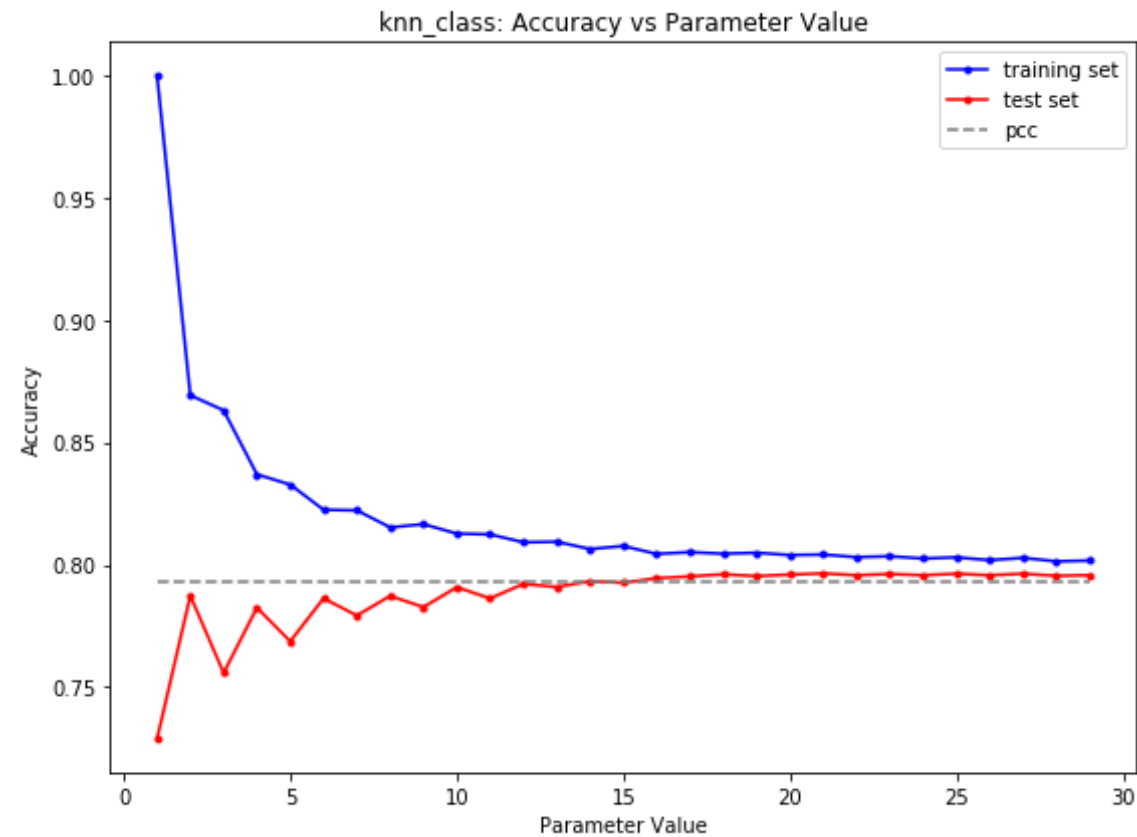
```
Population per class:      0
0  24720
1   7841
```

1.25 \* Proportion Chance Criterion: 79.29492137783143%



```
In [52]: ml_model(X, y, ml_type='knn_class', show_PCC=True,
               test_size=0.25, param_range=range(1, 30),
               seed_settings=range(0, 1), plot = True,
               report = True);
```





Report:

=====

Max accuracy: 0.7965

Optimal parameter: 21

1.25 x PCC: 0.7929

Total iterations: 1

In [ ]: